



Feature Normalization and Regularization

CEC Seminar

Johan Larsson

<https://jolars.co>, jola@math.ku.dk, [@jolars@mastodon.social](https://mstdn.social/@jolars)

Department of Mathematical Sciences, Copenhagen University

November 13, 2024

Who am I?

Problem and Motivation

Feature normalization has large effects in regularized regression (lasso, ridge) but there is no research on this.

Problem and Motivation

Feature normalization has large effects in regularized regression (lasso, ridge) but there is no research on this.

Results

- Class balance has a normalization-dependent impact on the model.
- In mixed data, choice of normalization implicitly weighs features' importances.



Figure 1: Joint work with Jonas Wallin

Preliminaries

Motivation and Aims

A Little Theory

Experiments

Preliminaries

- Data consists of a **fixed** matrix of features $\mathbf{X} \in \mathbb{R}^{n \times p}$ and a response vector $\mathbf{y} \in \mathbb{R}^n$.
- $\mathbf{y}$ comes from a linear model, that is,

$$y_i = \beta_0^* + \mathbf{x}_i^\top \boldsymbol{\beta}^* + \varepsilon_i \quad \text{for } i \in 1, \dots, n,$$

where $\boldsymbol{\beta}^*$ is the vector of *true* coefficients.

- ε_i is the measurement noise, generated from some random variable.

The Elastic Net

Linear regression plus a combination of the ℓ_1 and ℓ_2 penalties:

$$\beta^* = \arg \min_{\beta \in \mathbb{R}^p} \left(\frac{1}{2} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \underbrace{\lambda_1 \|\beta\|_1}_{\text{lasso}} + \underbrace{\frac{\lambda_2}{2} \|\beta\|_2^2}_{\text{ridge}} \right)$$

The Elastic Net

Linear regression plus a combination of the ℓ_1 and ℓ_2 penalties:

$$\beta^* = \arg \min_{\beta \in \mathbb{R}^p} \left(\frac{1}{2} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \underbrace{\lambda_1 \|\beta\|_1}_{\text{lasso}} + \underbrace{\frac{\lambda_2}{2} \|\beta\|_2^2}_{\text{ridge}} \right)$$




Figure 2: Elastic net is a combination of the lasso and ridge.

Regularization




Figure 3: Regularization for a simple linear regression problem




Figure 3: Regularization for a simple linear regression problem

Why Regularize? (Lasso, Ridge)

- Uniqueness when $p \gg n$
- To overcome overfitting




Figure 3: Regularization for a simple linear regression problem

Why Regularize? (Lasso, Ridge)

- Uniqueness when $p \gg n$
- To overcome overfitting

Why Sparsity? (Lasso)

- Interpretability
- The sparsity bet

- Don't know optimal λ_1 and λ_2 in advance.

The Elastic Net Path

- Don't know optimal λ_1 and λ_2 in advance.
- Instead we use hyper-parameter optimization (e.g. cross-validation).

- Don't know optimal λ_1 and λ_2 in advance.
- Instead we use hyper-parameter optimization (e.g. cross-validation).
- Common parametrization:

$$\begin{aligned}\lambda_1 &= \alpha\lambda, \\ \lambda_2 &= (1 - \alpha)\lambda\end{aligned}$$

with $\alpha \in [0, 1]$.

The Elastic Net Path

- Don't know optimal λ_1 and λ_2 in advance.
- Instead we use hyper-parameter optimization (e.g. cross-validation).
- Common parametrization:

$$\begin{aligned}\lambda_1 &= \alpha\lambda, \\ \lambda_2 &= (1 - \alpha)\lambda\end{aligned}$$

with $\alpha \in [0, 1]$.

- For each α , solve the elastic net over a sequence of λ : the **elastic net path**.




Figure 4: The elastic net path

Sensitivity to Scale

Lasso and ridge penalties are **norms**—feature scales matter!

Sensitivity to Scale

Lasso and ridge penalties are **norms**—feature scales matter!

Example

Assume

$$\mathbf{X} \sim \text{Normal} \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 4 & 0 \\ 0 & 1 \end{bmatrix} \right), \quad \boldsymbol{\beta}^* = \begin{bmatrix} \frac{1}{2} \\ 1 \end{bmatrix}.$$

and set $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \varepsilon$, with $\varepsilon \sim \text{Normal}(0, \sigma_\varepsilon^2)$.

Sensitivity to Scale

Lasso and ridge penalties are **norms**—feature scales matter!

Example

Assume

$$\mathbf{X} \sim \text{Normal} \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 4 & 0 \\ 0 & 1 \end{bmatrix} \right), \quad \boldsymbol{\beta}^* = \begin{bmatrix} \frac{1}{2} \\ 1 \end{bmatrix}.$$

and set $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \varepsilon$, with $\varepsilon \sim \text{Normal}(0, \sigma_\varepsilon^2)$.

Model	$\hat{\boldsymbol{\beta}}$	$\hat{\boldsymbol{\beta}}_{\text{std}}$
OLS	$[0.50 \quad 1.00]^\top$	$[1.00 \quad 1.00]^\top$

Sensitivity to Scale

Lasso and ridge penalties are **norms**—feature scales matter!

Example

Assume

$$\mathbf{X} \sim \text{Normal} \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 4 & 0 \\ 0 & 1 \end{bmatrix} \right), \quad \boldsymbol{\beta}^* = \begin{bmatrix} \frac{1}{2} \\ 1 \end{bmatrix}.$$

and set $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \varepsilon$, with $\varepsilon \sim \text{Normal}(0, \sigma_\varepsilon^2)$.

Model	$\hat{\boldsymbol{\beta}}$	$\hat{\boldsymbol{\beta}}_{\text{std}}$
OLS	$\begin{bmatrix} 0.50 & 1.00 \end{bmatrix}^\top$	$\begin{bmatrix} 1.00 & 1.00 \end{bmatrix}^\top$
Lasso	$\begin{bmatrix} 0.38 & 0.50 \end{bmatrix}^\top$	$\begin{bmatrix} 0.74 & 0.50 \end{bmatrix}^\top$

Sensitivity to Scale

Lasso and ridge penalties are **norms**—feature scales matter!

Example

Assume

$$\mathbf{X} \sim \text{Normal} \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 4 & 0 \\ 0 & 1 \end{bmatrix} \right), \quad \boldsymbol{\beta}^* = \begin{bmatrix} \frac{1}{2} \\ 1 \end{bmatrix}.$$

and set $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \varepsilon$, with $\varepsilon \sim \text{Normal}(0, \sigma_\varepsilon^2)$.

Model	$\hat{\boldsymbol{\beta}}$	$\hat{\boldsymbol{\beta}}_{\text{std}}$
OLS	$\begin{bmatrix} 0.50 & 1.00 \end{bmatrix}^\top$	$\begin{bmatrix} 1.00 & 1.00 \end{bmatrix}^\top$
Lasso	$\begin{bmatrix} 0.38 & 0.50 \end{bmatrix}^\top$	$\begin{bmatrix} 0.74 & 0.50 \end{bmatrix}^\top$
Ridge	$\begin{bmatrix} 0.37 & 0.41 \end{bmatrix}^\top$	$\begin{bmatrix} 0.74 & 0.41 \end{bmatrix}^\top$

Sensitivity to Scale

Lasso and ridge penalties are **norms**—feature scales matter!

Example

Assume

$$\mathbf{X} \sim \text{Normal} \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 4 & 0 \\ 0 & 1 \end{bmatrix} \right), \quad \boldsymbol{\beta}^* = \begin{bmatrix} \frac{1}{2} \\ 1 \end{bmatrix}.$$

and set $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \varepsilon$, with $\varepsilon \sim \text{Normal}(0, \sigma_\varepsilon^2)$.

Model	$\hat{\boldsymbol{\beta}}$	$\hat{\boldsymbol{\beta}}_{\text{std}}$
OLS	$\begin{bmatrix} 0.50 & 1.00 \end{bmatrix}^\top$	$\begin{bmatrix} 1.00 & 1.00 \end{bmatrix}^\top$
Lasso	$\begin{bmatrix} 0.38 & 0.50 \end{bmatrix}^\top$	$\begin{bmatrix} 0.74 & 0.50 \end{bmatrix}^\top$
Ridge	$\begin{bmatrix} 0.37 & 0.41 \end{bmatrix}^\top$	$\begin{bmatrix} 0.74 & 0.41 \end{bmatrix}^\top$

Large scale means less penalization because the size of β_j can be smaller for an equivalent effect (on $\mathbf{y}$).

Scale sensitivity can be mitigated by normalizing the features. Let $\tilde{X}$ be the normalized feature matrix, with elements

$$\tilde{x}_{ij} = \frac{x_{ij} - c_j}{s_j}.$$

Scale sensitivity can be mitigated by normalizing the features. Let $\tilde{\mathbf{X}}$ be the normalized feature matrix, with elements

$$\tilde{x}_{ij} = \frac{x_{ij} - c_j}{s_j}.$$

After fitting, we transform the coefficients back to their original scale via

$$\hat{\beta}_j = \frac{\hat{\beta}_j^{(n)}}{s_j} \quad \text{for } j = 1, 2, \dots, p,$$

where $\hat{\beta}_j^{(n)}$ is a coefficient from the normalized problem.

Table 1: Common ways to normalize X

Normalization	Centering (c_j)	Scaling (s_j)
Standardization	$\bar{x}_j = \frac{1}{n} \sum_{i=1}^n x_{ij}$	$\sqrt{\frac{1}{n} \sum_{i=1}^n (x_{ij} - \bar{x}_j)^2}$
Min–Max	$\min_i(x_{ij})$	$\max_i(x_{ij}) - \min_i(x_{ij})$
Max–Abs	0	$\max_i(x_{ij})$
Adaptive Lasso	0	β_j^{OLS}

Motivation and Aims




Figure 5: Normalization matters. Lasso paths under two different types of normalization (standardization and max-abs normalization). The union of the first ten features selected in any of the settings are colored.

Table 2: Lasso coefficients on test sets, with λ set from 5-fold cross-validation repeated 5 times. We show the five largest coefficients in magnitude for each data set for the standardization setting and corresponding coefficients for the max-abs normalization setting.

housing		triazines		w1a	
$\hat{\beta}_{\text{std}}$	$\hat{\beta}_{\text{max-abs}}$	$\hat{\beta}_{\text{std}}$	$\hat{\beta}_{\text{max-abs}}$	$\hat{\beta}_{\text{std}}$	$\hat{\beta}_{\text{max-abs}}$
-0.63	-0.68	0.17	0.0	1.8	0.0
-1.4	-0.78	0.069	0.0	1.8	0.78
0.27	0.0	0.028	0.0	1.8	0.63
-0.99	-0.34	0.071	0.0	1.4	0.080
2.8	3.1	0.029	0.0	1.7	0.0

Motivation

Normalization matters but there is no research into this.

Motivation

Normalization matters but there is no research into this.

Everyone agrees you need to normalize, but how to do so is usually motivated by being “standard”.

Motivation

Normalization matters but there is no research into this.

Everyone agrees you need to normalize, but how to do so is usually motivated by being “standard”.

The meaning of “standard” depends on field!

Motivation

Normalization matters but there is no research into this.

Everyone agrees you need to normalize, but how to do so is usually motivated by being “standard”.

The meaning of “standard” depends on field!

What We Show

Normalization influences shrinkage for the elastic net for binary features.

Motivation

Normalization matters but there is no research into this.

Everyone agrees you need to normalize, but how to do so is usually motivated by being “standard”.

The meaning of “standard” depends on field!

What We Show

Normalization influences shrinkage for the elastic net for binary features.

There's a bias–variance tradeoff that depends on the normalization.

Motivation

Normalization matters but there is no research into this.

Everyone agrees you need to normalize, but how to do so is usually motivated by being “standard”.

The meaning of “standard” depends on field!

What We Show

Normalization influences shrinkage for the elastic net for binary features.

There's a bias–variance tradeoff that depends on the normalization.

For the elastic net, you should scale the penalty weights—not the features.

Motivation

Normalization matters but there is no research into this.

Everyone agrees you need to normalize, but how to do so is usually motivated by being “standard”.

The meaning of “standard” depends on field!

What We Show

Normalization influences shrinkage for the elastic net for binary features.

There's a bias–variance tradeoff that depends on the normalization.

For the elastic net, you should scale the penalty weights—not the features.

When mixing binary and normal features, normalization implicitly weighs their importance.

A Little Theory

Orthogonal Features

There is no explicit solution to the elastic net problem in general.

¹We have also assumed that the features are mean-centered here.

Orthogonal Features

There is no explicit solution to the elastic net problem in general.

But if we assume that the features are orthogonal, that is

$$\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} = \text{diag}(\tilde{\mathbf{x}}_1^\top \tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_p^\top \tilde{\mathbf{x}}_p),$$

then there is¹:

$$\hat{\beta}_j = \frac{S_{\lambda_1}(\tilde{\mathbf{x}}_j^\top \mathbf{y})}{s_j(\tilde{\mathbf{x}}_j^\top \tilde{\mathbf{x}}_j + \lambda_2)},$$

where

$$S_\lambda(z) = \text{sign}(z) \max(|z| - \lambda, 0).$$

¹We have also assumed that the features are mean-centered here.

Orthogonal Features

There is no explicit solution to the elastic net problem in general.

But if we assume that the features are orthogonal, that is

$$\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} = \text{diag}(\tilde{\mathbf{x}}_1^\top \tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_p^\top \tilde{\mathbf{x}}_p),$$

then there is¹:

$$\hat{\beta}_j = \frac{S_{\lambda_1}(\tilde{\mathbf{x}}_j^\top \mathbf{y})}{s_j(\tilde{\mathbf{x}}_j^\top \tilde{\mathbf{x}}_j + \lambda_2)},$$

where

$$S_\lambda(z) = \text{sign}(z) \max(|z| - \lambda, 0).$$




Figure 6: Soft thresholding

¹We have also assumed that the features are mean-centered here.

Binary Features

Assume we have a binary feature x_j , such that $x_{ij} \in \{0, 1\}$. Let $q \in [0, 1]$ be the class balance of this feature, that is: $q = \bar{x}_j$.

Binary Features

Assume we have a binary feature x_j , such that $x_{ij} \in \{0, 1\}$. Let $q \in [0, 1]$ be the class balance of this feature, that is: $q = \bar{x}_j$.

Noiseless Case

In the noiseless case, we have

$$\hat{\beta}_j = \frac{S_{\lambda_1} \left(\frac{\beta_j^* n(q - q^2)}{s_j} \right)}{s_j \left(\frac{n(q - q^2)}{s_j^2} + \lambda_2 \right)}.$$

Binary Features

Assume we have a binary feature x_j , such that $x_{ij} \in \{0, 1\}$. Let $q \in [0, 1]$ be the class balance of this feature, that is: $q = \bar{x}_j$.

Noiseless Case

In the noiseless case, we have

$$\hat{\beta}_j = \frac{S_{\lambda_1} \left(\frac{\beta_j^* n(q - q^2)}{s_j} \right)}{s_j \left(\frac{n(q - q^2)}{s_j^2} + \lambda_2 \right)}.$$

Means that the elastic net estimator depends on class balance (q).

Binary Features

Assume we have a binary feature x_j , such that $x_{ij} \in \{0, 1\}$. Let $q \in [0, 1]$ be the class balance of this feature, that is: $q = \bar{x}_j$.

Noiseless Case

In the noiseless case, we have

$$\hat{\beta}_j = \frac{S_{\lambda_1} \left(\frac{\beta_j^* n(q - q^2)}{s_j} \right)}{s_j \left(\frac{n(q - q^2)}{s_j^2} + \lambda_2 \right)}.$$

Means that the elastic net estimator depends on class balance (q).

$s_j = q - q^2$ for lasso and $s_j = \sqrt{q - q^2}$ for ridge removes effect of q , which suggests the parametrization

$$s_j = (q - q^2)^\delta, \quad \delta \geq 0.$$

Binary Features

Assume we have a binary feature x_j , such that $x_{ij} \in \{0, 1\}$. Let $q \in [0, 1]$ be the class balance of this feature, that is: $q = \bar{x}_j$.

Noiseless Case

In the noiseless case, we have

$$\hat{\beta}_j = \frac{S_{\lambda_1} \left(\frac{\beta_j^* n(q - q^2)}{s_j} \right)}{s_j \left(\frac{n(q - q^2)}{s_j^2} + \lambda_2 \right)}.$$

Means that the elastic net estimator depends on class balance (q).

$s_j = q - q^2$ for lasso and $s_j = \sqrt{q - q^2}$ for ridge removes effect of q , which suggests the parametrization

$$s_j = (q - q^2)^\delta, \quad \delta \geq 0.$$

There is no (simple) s_j that will work for the elastic net.

Probability of Selection

Since $\mathbf{X}$ is fixed and ε is normal, we can compute the probability of selection:

$$\Pr(\hat{\beta}_j \neq 0) = \Phi\left(\frac{\mu - \lambda_1}{\sigma}\right) + \Phi\left(\frac{-\mu - \lambda_1}{\sigma}\right).$$




Figure 7: Probability that the elastic net selects a feature across different noise levels (σ_ε), types of normalization (δ), and class balance (q). The dashed line is asymptotic behavior for $\delta = 1/2$. Scaling used is $s_j \propto (q - q^2)^\delta$.

Rare Traits

Features with large class-imbalances might not be selected even if effect is **very strong** (e.g. rare SNPs, mutations).

Rare Traits

Features with large class-imbalances might not be selected even if effect is **very strong** (e.g. rare SNPs, mutations).

Subgroup Data

Results become dependent on data collection.

Collecting more data with different class balances influences the results (since class balances change).

Theorem

If x_j is a binary feature with class balance $q \in (0, 1)$ and $\lambda_1, \lambda_2 \in (0, \infty)$, $\sigma_\epsilon > 0$, and $s_j = (q - q^2)^\delta$, $\delta \geq 0$, then

$$\lim_{q \rightarrow 1^+} \mathbb{E} \hat{\beta}_j = \begin{cases} 0 & \text{if } 0 \leq \delta < \frac{1}{2}, \\ \frac{2n\beta_j^*}{n+\lambda_2} \square \left(-\frac{\lambda_1}{\sigma_\epsilon \sqrt{n}} \right) & \text{if } \delta = \frac{1}{2}, \\ \beta_j^* & \text{if } \delta \geq \frac{1}{2}. \end{cases}$$

Theorem

If x_j is a binary feature with class balance $q \in (0, 1)$ and $\lambda_1, \lambda_2 \in (0, \infty)$, $\sigma_\epsilon > 0$, and $s_j = (q - q^2)^\delta$, $\delta \geq 0$, then

$$\lim_{q \rightarrow 1^+} \mathbb{E} \hat{\beta}_j = \begin{cases} 0 & \text{if } 0 \leq \delta < \frac{1}{2}, \\ \frac{2n\beta_j^*}{n+\lambda_2} \square \left(-\frac{\lambda_1}{\sigma_\epsilon \sqrt{n}} \right) & \text{if } \delta = \frac{1}{2}, \\ \beta_j^* & \text{if } \delta \geq \frac{1}{2}. \end{cases}$$

and

$$\lim_{q \rightarrow 1^+} \text{Var} \hat{\beta}_j = \begin{cases} 0 & \text{if } 0 \leq \delta < \frac{1}{2}, \\ \infty & \text{if } \delta \geq \frac{1}{2}. \end{cases}$$






Figure 8: A bias variance tradeoff. Bias, variance, and mean-squared error for a one-dimensional lasso problem. Theoretical result for orthogonal features. Dotted line is asymptotic result or $\delta = 1/2$. Scaling used is $s_j \propto (q - q^2)^\delta$.




(a) Power in the sense of detecting all the true signals. Constant p .

Figure 9: Multiple features: 10 true signals and varying q and p . Mean squared error (MSE), false discovery rate (FDR), and power



(a) Power in the sense of detecting all the true signals. Constant p .



(b) False discovery rate (FDR) and normalized mean-squared error (NMSE).

Figure 9: Multiple features: 10 true signals and varying q and p . Mean squared error (MSE), false discovery rate (FDR), and power

So far: all binary features. What about mixing binary and continuous (normal) features?

How to put binary features and normal features on the “same” scale?




Figure 10: How do we match these?

Our Definition of Comparability

The effects of a binary feature and a normally distributed feature are **comparable** if a flip in the binary feature has the same effect as a two-standard deviation change in the normal feature (Gelman 2008).

Our Definition of Comparability

The effects of a binary feature and a normally distributed feature are **comparable** if a flip in the binary feature has the same effect as a two-standard deviation change in the normal feature (Gelman 2008).

Examples

Assume entries in x_1 are binary and x_2 come from a random variable X_2 . The effects are comparable in the following cases:

- $X_2 \sim \text{Normal}(\mu, 1/2)$, $\beta_1^* = 1$, and $\beta_2^* = 1$.
- $X_2 \sim \text{Normal}(\mu, 2)$, $\beta_1^* = 1$, and $\beta_2^* = 0.25$.

For the two-standard deviation notion of comparability to hold, we need to modify our scaling factor s_j .

Choice of Scaling in Mixed Data

For the two-standard deviation notion of comparability to hold, we need to modify our scaling factor s_j .

As before, we assume that x_1 is binary and $X_2 \sim \text{Normal}(\mu, 1/2)$, $\beta_1^* = \beta_2^* = 1$ so that they have *comparable* effects. Also assume we standardize x_2 .

We want $\hat{\beta}_1 = \hat{\beta}_2$. That is,

$$\underbrace{\frac{S_{\lambda_1} \left(\frac{n(q-q^2)}{s_j} \right)}{s_1 \left(\frac{n(q-q^2)}{s_1^2} + \lambda_2 \right)}}_{\hat{\beta}_1} = \underbrace{\frac{S_{\lambda_1} \left(\frac{n}{2} \right)}{\frac{1}{2} (n + \lambda_2)}}_{\hat{\beta}_2}.$$

Choice of Scaling in Mixed Data

For the two-standard deviation notion of comparability to hold, we need to modify our scaling factor s_j .

As before, we assume that x_1 is binary and $X_2 \sim \text{Normal}(\mu, 1/2)$, $\beta_1^* = \beta_2^* = 1$ so that they have *comparable* effects. Also assume we standardize x_2 .

We want $\hat{\beta}_1 = \hat{\beta}_2$. That is,

$$\underbrace{\frac{S_{\lambda_1} \left(\frac{n(q-q^2)}{s_j} \right)}{s_1 \left(\frac{n(q-q^2)}{s_1^2} + \lambda_2 \right)}}_{\hat{\beta}_1} = \underbrace{\frac{S_{\lambda_1} \left(\frac{n}{2} \right)}{\frac{1}{2} (n + \lambda_2)}}_{\hat{\beta}_2}.$$

The choice $s_1 = (2(q - q^2))^\delta$ works when classes are balanced ($q = 0.5$). But no clear choice for the elastic net case.

Experiments

Binary Features (Decreasing q)




Figure 11: Lasso estimates for first 30 coefficients. First 20 features are true signals with a geometrically decreasing class balance from 0.5 to 0.99. ρ is a measure of autocorrelation.

In the middle

Binary Features (Signal-to-Noise Ratio)




Figure 12: Normalized mean-squared test set error (NMSE).

Mixed Data




Figure 13: Comparison between lasso and ridge estimators for features generated to resemble features from various distributions.

Hyperparameter Optimization




Figure 14: Contour plots of hold-out (validation set) error across a grid of δ and λ values for the lasso and ridge.

Conclusions

- Class balance plays a crucial role when using regularized regression on binary data.
- As far as we know the first paper to investigate the interplay between normalization and regularization
- New scaling approach to deal with class-imbalanced binary features
- Discussion and suggestions for dealing with mixed data

Conclusions

- Class balance plays a crucial role when using regularized regression on binary data.
- As far as we know the first paper to investigate the interplay between normalization and regularization
- New scaling approach to deal with class-imbalanced binary features
- Discussion and suggestions for dealing with mixed data

Limitations

- So far only theoretical results for limited cases:
 - Fixed data (X), normal noise
 - Orthogonal features
 - Normal and binary features

Thank you!

-  El Ghaoui, Laurent, Vivian Viallon, and Tarek Rabbani (Sept. 21, 2010). *Safe Feature Elimination in Sparse Supervised Learning*. Technical report UCB/EECS-2010-126. Berkeley: EECS Department, University of California. URL: <http://www.eecs.berkeley.edu/Pubs/TechRpts/2010/EECS-2010-126.html>.
-  Gelman, Andrew (July 10, 2008). “Scaling Regression Inputs by Dividing by Two Standard Deviations”. In: *Statistics in Medicine* 27.15, pp. 2865–2873. ISSN: 02776715, 10970258. DOI: [10.1002/sim.3107](https://doi.org/10.1002/sim.3107). URL: <https://onlinelibrary.wiley.com/doi/10.1002/sim.3107> (visited on 09/27/2023).
-  Tibshirani, Robert et al. (Mar. 2012). “Strong Rules for Discarding Predictors in Lasso-Type Problems”. In: *Journal of the Royal Statistical Society. Series B: Statistical Methodology* 74.2, pp. 245–266. ISSN: 1369-7412. DOI: [10/c4bb85](https://doi.org/10.1016/j.jrssb.2012.03.001). URL: <https://iths.pure.elsevier.com/en/publications/strong-rules-for-discarding-predictors-in-lasso-type-problems> (visited on 03/16/2018).

-  Zou, Hui and Trevor Hastie (2005). “Regularization and Variable Selection via the Elastic Net”. In: *Journal of the Royal Statistical Society. Series B (Statistical Methodology)* 67.2, pp. 301–320. ISSN: 1369-7412. URL: www.jstor.org/stable/3647580 (visited on 03/12/2018).

Extras

Max-Abs Scaling of Continuous Features

- Min-max normalization is sometimes used in continuous data
- Very sensitive to outliers
- But also depend on sample size!
- In other words, results in model validation with varying sample sizes can yield very strange results.




Figure 15: Effects of maximum absolute value scaling.

Hyperparameter Optimization (Support and NMSE)




Figure 16: Support and NMSE of the lasso for different values of δ and λ .

- Proposed by Zou and Hastie ([2005](#)).

- Proposed by Zou and Hastie (2005).
- Lasso (ℓ_1) part:
 - Enables sparsity (interpretability, parsimony, feature selection)
 - Efficient when $p \gg n$ (due to screening rules (El Ghaoui, Viallon, and Rabbani 2010; Tibshirani et al. 2012))

- Proposed by Zou and Hastie (2005).
- Lasso (ℓ_1) part:
 - Enables sparsity (interpretability, parsimony, feature selection)
 - Efficient when $p \gg n$ (due to screening rules (El Ghaoui, Viallon, and Rabbani 2010; Tibshirani et al. 2012))
- Ridge (ℓ_2) part
 - Mitigates lasso issue in correlated data
 - Better predictive performance when true signal is non-sparse